

# 从零开始的RISC-V模拟器开发

## 第4讲 Spike篇之内存模拟

中国科学院软件研究所  
PLCT实验室

李威威 liweiwei@iscas.ac.cn

王俊强 wangjunqiang@iscas.ac.cn

吴伟 wuwei2016@iscas.ac.cn

# 回顾

- 指令执行流程: riscv/execute.cc: step()

```
auto ic_entry = _mmu->access_icache(pc);

#define ICACHE_ACCESS(i) { \
    insn_fetch_t fetch = ic_entry->data; \
    pc = execute_insn(this, pc, fetch); \
    ic_entry = ic_entry->next; \
    if (i == mmu_t::ICACHE_ENTRIES-1) break; \
    if (unlikely(ic_entry->tag != pc)) break; \
    if (unlikely(instret+1 == n)) break; \
    instret++; \
    state.pc = pc; \
}
```

```
static reg_t execute_insn(processor_t* p, reg_t pc,
    insn_fetch_t fetch)
{
    ...
    npc = fetch.func(p, fetch.insn, pc);
    ...
}
```

- riscv/mmu.h

```
struct insn_fetch_t
{
    insn_func_t func;
    insn_t insn;
};
```

```
inline icache_entry_t* access_icache(reg_t addr)
{
    icache_entry_t* entry = &icache[icache_index(addr)];
    if (likely(entry->tag == addr))
        return entry;
    return refill_icache(addr, entry);
}
```


```
inline icache_entry_t* refill_icache(reg_t addr, icache_entry_t* entry)
{
    auto tlb_entry = translate_insn_addr(addr);
    insn_bits_t insn = from_le(*(uint16_t*)(tlb_entry.host_offset + addr));
    int length = insn_length(insn);
    ...
    insn_fetch_t fetch = {proc->decode_insn(insn), insn};
    entry->tag = addr;
    entry->next = &icache[icache_index(addr + length)];
    entry->data = fetch;
    ...
    return entry;
}
```

# Spike内存系统层次结构

spike\_main/spike.cc

```
std::unique_ptr<icache_sim_t> ic;  
std::unique_ptr<dcache_sim_t> dc;  
std::unique_ptr<cache_sim_t> l2;  
...  
sim_t s(isa, priv, varch, nprocs, halted, real_time_clint, initrd_start,  
initrd_end, bootargs, start_pc, mems, plugin_devices, htif_args,  
std::move(hartids), dm_config, log_path, dtb_enabled, dtb_file);  
...  
if (ic && l2) ic->set_miss_handler(&*l2);  
if (dc && l2) dc->set_miss_handler(&*l2);  
...  
for (size_t i = 0; i < nprocs; i++)  
{  
    if (ic) s.get_core(i)->get_mmu()->register_memtracer(&*ic);  
    if (dc) s.get_core(i)->get_mmu()->register_memtracer(&*dc);  
    ...  
}
```

```
for (auto& x : mems)  
    bus.add_device(x.first, x.second);
```




# Spike内存系统—IC和DC

riscv/cachesim.h

```
class icache_sim_t : public cache_memtracer_t
{
public:
    icache_sim_t(const char* config) : cache_memtracer_t(config, "I$") {}
    bool interested_in_range(uint64_t begin, uint64_t end, access_type type)
    {
        return type == FETCH;
    }
    void trace(uint64_t addr, size_t bytes, access_type type)
    {
        if (type == FETCH) cache->access(addr, bytes, false);
    }
};
```

```
class dcache_sim_t : public cache_memtracer_t
{
public:
    dcache_sim_t(const char* config) :
        cache_memtracer_t(config, "D$") {}
    bool interested_in_range(uint64_t begin, uint64_t end,
        access_type type)
    {
        return type == LOAD || type == STORE;
    }
    void trace(uint64_t addr, size_t bytes, access_type type)
    {
        if (type == LOAD || type == STORE) cache->access(addr,
            bytes, type == STORE);
    }
};
```




# Spike内存系统—cache\_memtracer\_t

riscv/cachesim.h

```
class cache_memtracer_t : public memtracer_t
{
public:
    cache_memtracer_t(const char* config, const char* name)
    {
        cache = cache_sim_t::construct(config, name);
    }
    ~cache_memtracer_t()
    {
        delete cache;
    }
    void set_miss_handler(cache_sim_t* mh)
    {
        cache->set_miss_handler(mh);
    }
    void set_log(bool log)
    {
        cache->set_log(log);
    }
protected:
    cache_sim_t* cache;
};
```

```
class memtracer_t
{
public:
    memtracer_t() {}
    virtual ~memtracer_t() {}
    virtual bool interested_in_range(uint64_t begin, uint64_t end, access_type
type) = 0;
    virtual void trace(uint64_t addr, size_t bytes, access_type type) = 0;
};
```



riscv/memtracer.h

# Spike内存系统—cache\_sim\_t及L2

riscv/cachesim.h

```
class cache_sim_t
{
public:
    cache_sim_t(size_t sets, size_t ways, size_t linesz, const char* name);
    cache_sim_t(const cache_sim_t& rhs);
    virtual ~cache_sim_t();

    void access(uint64_t addr, size_t bytes, bool store);
    void print_stats();
    void set_miss_handler(cache_sim_t* mh) { miss_handler = mh; }
    void set_log(bool _log) { log = _log; }

    static cache_sim_t* construct(const char* config, const char* name);


protected:
    static const uint64_t VALID = 1ULL << 63;
    static const uint64_t DIRTY = 1ULL << 62;

    virtual uint64_t* check_tag(uint64_t addr);
    virtual uint64_t victimize(uint64_t addr);
```

```
lfsr_t lfsr;
cache_sim_t* miss_handler;
size_t sets;
size_t ways;
size_t linesz;
size_t idx_shift;
uint64_t* tags;
uint64_t read_accesses;
uint64_t read_misses;
uint64_t bytes_read;
uint64_t write_accesses;
uint64_t write_misses;
uint64_t bytes_written;
uint64_t writebacks;

std::string name;
bool log;

void init();
};
```



# Spike内存系统—cache access

riscv/cachesim.cc

```
void cache_sim_t::access(uint64_t addr, size_t bytes, bool store)
{
    store ? write_accesses++ : read_accesses++;
    (store ? bytes_written : bytes_read) += bytes;

    uint64_t* hit_way = check_tag(addr);
    if (likely(hit_way != NULL))
    {
        if (store)
            *hit_way |= DIRTY;
        return;
    }


    store ? write_misses++ : read_misses++;
    if (log)
    {
        std::cerr << name << " "
            << (store ? "write" : "read") << " miss 0x"
            << std::hex << addr << std::endl;
    }
}
```

```
uint64_t victim = victimize(addr);

if ((victim & (VALID | DIRTY)) == (VALID | DIRTY))
{
    uint64_t dirty_addr = (victim & ~(VALID | DIRTY)) << idx_shift;
    if (miss_handler)
        miss_handler->access(dirty_addr, linesz, true);
    writebacks++;
}

if (miss_handler)
    miss_handler->access(addr & ~(linesz-1), linesz, false);

if (store)
    *check_tag(addr) |= DIRTY;
}
```






# Spike内存系统—MMU

riscv/mmu.h

```
class mmu_t
{
private:
    std::map<reg_t, reg_t> alloc_cache;
    std::vector<std::pair<reg_t, reg_t>> addr_tbl;
...
    memtracer_list_t tracer;
...
    // implement an instruction cache for simulator performance
    icache_entry_t icache[ICACHE_ENTRIES];
...
    void register_memtracer(memtracer_t*);
...
}
```

```
void mmu_t::register_memtracer(memtracer_t* t)
{
    flush_tlb();
    tracer.hook(t);
}
```






# Spike内存系统—memtracer\_list\_t

riscv/mmu.h

```
class memtracer_list_t : public memtracer_t
{
public:
    bool empty() { return list.empty(); }
    bool interested_in_range(uint64_t begin, uint64_t end, access_type type)
    {
        for (std::vector<memtracer_t*>::iterator it = list.begin(); it != list.end(); ++it)
            if ((*it)->interested_in_range(begin, end, type))
                return true;
        return false;
    }
    void trace(uint64_t addr, size_t bytes, access_type type)
    {
        for (std::vector<memtracer_t*>::iterator it = list.begin(); it != list.end(); ++it)
            (*it)->trace(addr, bytes, type);
    }
    void hook(memtracer_t* h)
    {
        list.push_back(h);
    }
private:
    std::vector<memtracer_t*> list;
};
```



# 快速指令执行流程

- 指令执行流程: riscv/execute.cc: step()

```
auto ic_entry = _mmu->access_icache(pc);

#define ICACHE_ACCESS(i) { \
    insn_fetch_t fetch = ic_entry->data; \
    pc = execute_insn(this, pc, fetch); \
    ic_entry = ic_entry->next; \
    if (i == mmu_t::ICACHE_ENTRIES-1) break; \
    if (unlikely(ic_entry->tag != pc)) break; \
    if (unlikely(instret+1 == n)) break; \
    instret++; \
    state.pc = pc; \
}
```

```
static reg_t execute_insn(processor_t* p, reg_t pc,
    insn_fetch_t fetch)
{
    ...
    npc = fetch.func(p, fetch.insn, pc);
    ...
}
```

- riscv/mmu.h

```
struct insn_fetch_t
{
    insn_func_t func;
    insn_t insn;
};
```

```
inline icache_entry_t* access_icache(reg_t addr)
{
    icache_entry_t* entry = &icache[icache_index(addr)];
    if (likely(entry->tag == addr))
        return entry;
    return refill_icache(addr, entry);
}
```

```
inline icache_entry_t* refill_icache(reg_t addr, icache_entry_t* entry)
{
    auto tlb_entry = translate_insn_addr(addr);
    insn_bits_t insn = from_le(*(uint16_t*)(tlb_entry.host_offset + addr));
    int length = insn_length(insn);
    ...
    insn_fetch_t fetch = {proc->decode_insn(insn), insn};
    entry->tag = addr;
    entry->next = &icache[icache_index(addr + length)];
    entry->data = fetch;
    ...
    return entry;
}
```

# 慢速指令执行流程




- 指令执行流程: riscv/execute.cc: step()

```
insn_fetch_t fetch = mmu->load_insn(pc);  
if (debug && !state.serialized)  
    disasm(fetch.insn);  
pc = execute_insn(this, pc, fetch);  
advance_pc();
```



- riscv/mmu.h

```
inline insn_fetch_t load_insn(reg_t addr)  
{  
    icache_entry_t entry;  
    return refill_icache(addr, &entry)->data;  
}
```



```
struct insn_fetch_t  
{  
    insn_func_t func;  
    insn_t insn;  
};
```



```
static reg_t execute_insn(processor_t* p, reg_t pc,  
    insn_fetch_t fetch)  
{  
    ...  
    npc = fetch.func(p, fetch.insn, pc);  
    ...  
}
```


```
inline icache_entry_t* refill_icache(reg_t addr, icache_entry_t* entry)  
{  
    auto tlb_entry = translate_insn_addr(addr);  
    insn_bits_t insn = from_le(*(uint16_t*)(tlb_entry.host_offset + addr));  
    int length = insn_length(insn);  
    ...  
  
    insn_fetch_t fetch = {proc->decode_insn(insn), insn};  
    entry->tag = addr;  
    entry->next = &icache[icache_index(addr + length)];  
    entry->data = fetch;  
  
    ...  
    return entry;  
}
```

# Spike内存系统—取指流程

riscv/mmu.h

```
inline icache_entry_t* refill_icache(reg_t addr, icache_entry_t* entry)
{
    auto tlb_entry = translate_insn_addr(addr);
    insn_bits_t insn = from_le(*(uint16_t*)(tlb_entry.host_offset +
addr));
    ...
    insn_fetch_t fetch = {proc->decode_insn(insn), insn};
    entry->tag = addr;
    entry->next = &icache[icache_index(addr + length)];
    entry->data = fetch;


    reg_t paddr = tlb_entry.target_offset + addr;
    if (tracer.interested_in_range(paddr, paddr + 1, FETCH)) {
        entry->tag = -1;
        tracer.trace(paddr, length, FETCH);
    }
    return entry;
}
```



# Spike内存系统—地址转换

riscv/mmu.h

```
inline tlb_entry_t translate_insn_addr(reg_t addr) {  
    reg_t vpn = addr >> PGSHIFT;  
    if (likely(tlb_insn_tag[vpn % TLB_ENTRIES] == vpn))  
        return tlb_data[vpn % TLB_ENTRIES];  
    tlb_entry_t result;  
    if (unlikely(tlb_insn_tag[vpn % TLB_ENTRIES] != (vpn |  
    TLB_CHECK_TRIGGERS))) {  
        result = fetch_slow_path(addr);  
    } else {  
        result = tlb_data[vpn % TLB_ENTRIES];  
    }  
    if (unlikely(tlb_insn_tag[vpn % TLB_ENTRIES] == (vpn |  
    TLB_CHECK_TRIGGERS))) {  
        uint16_t* ptr = (uint16_t*)(tlb_data[vpn % TLB_ENTRIES].host_offset + addr);  
        int match = proc->trigger_match(OPERATION_EXECUTE, addr, from_le(*ptr));  
        if (match >= 0) {  
            throw trigger_matched_t(match, OPERATION_EXECUTE, addr, from_le(*ptr));  
        }  
    }  
    return result;  
}
```




# Spike内存系统扩展—取指转换慢速路径

riscv/mmu.cc

```
tlb_entry_t mmu_t::fetch_slow_path(reg_t vaddr)
{
    reg_t paddr = translate(vaddr, sizeof(fetch_temp), FETCH, 0);

    if (auto host_addr = sim->addr_to_mem(paddr)) {
        return refill_tlb(vaddr, paddr, host_addr, FETCH);
    } else {
        if (!mmio_load(paddr, sizeof fetch_temp, (uint8_t*)&fetch_temp))
            throw trap_instruction_access_fault(vaddr, 0, 0);
        tlb_entry_t entry = {(char*)&fetch_temp - vaddr, paddr - vaddr};
        return entry;
    }
}
```



# Spike内存系统扩展—取指转换慢速路径


riscv/mmu.cc

```
tlb_entry_t mmu_t::fetch_slow_path(reg_t vaddr)
{
    reg_t paddr = translate(vaddr, sizeof(fetch_temp), FETCH, 0);

    if (auto host_addr = sim->addr_to_mem(paddr)) {
        return refill_tlb(vaddr, paddr, host_addr, FETCH);
    } else {
        if (!mmio_load(paddr, sizeof fetch_temp, (uint8_t*)&fetch_temp))
            throw trap_instruction_access_fault(vaddr, 0, 0);
        tlb_entry_t entry = {(char*)&fetch_temp - vaddr, paddr - vaddr};
        return entry;
    }
}
```

riscv/sim.cc

```
char* sim_t::addr_to_mem(reg_t addr) {
    if (!paddr_ok(addr))
        return NULL;
    auto desc = bus.find_device(addr);
    if (auto mem = dynamic_cast<mem_t*>(desc.second))
        if (addr - desc.first < mem->size())
            return mem->contents() + (addr - desc.first);
    return NULL;
}
```






# Spike内存系统扩展—填充tlb

```
tlb_entry_t mmu_t::refill_tlb(reg_t vaddr, reg_t paddr, char* host_addr, access_type type)
{
    reg_t idx = (vaddr >> PGSHIFT) % TLB_ENTRIES;
    reg_t expected_tag = vaddr >> PGSHIFT;
    if ((tlb_load_tag[idx] & ~TLB_CHECK_TRIGGERS) != expected_tag)
        tlb_load_tag[idx] = -1;
    if ((tlb_store_tag[idx] & ~TLB_CHECK_TRIGGERS) != expected_tag)
        tlb_store_tag[idx] = -1;
    if ((tlb_insn_tag[idx] & ~TLB_CHECK_TRIGGERS) != expected_tag)
        tlb_insn_tag[idx] = -1;

    if ((check_triggers_fetch && type == FETCH) ||
        (check_triggers_load && type == LOAD) ||
        (check_triggers_store && type == STORE))
        expected_tag |= TLB_CHECK_TRIGGERS;

    if (pmp_homogeneous(paddr & ~reg_t(PGSIZE - 1), PGSIZE)) {
        if (type == FETCH) tlb_insn_tag[idx] = expected_tag;
        else if (type == STORE) tlb_store_tag[idx] = expected_tag;
        else tlb_load_tag[idx] = expected_tag;
    }
    tlb_entry_t entry = {host_addr - vaddr, paddr - vaddr};
    tlb_data[idx] = entry;
    return entry;
}
```



# Spike内存系统—加载内存


```
WRITE_RVC_RS2S(MMU.load_int32(RVC_RS1S + insn.rvc_lw_imm()));
```

```
load_func(int16, load, 0)  
load_func(int32, load, 0)
```

```
#define load_func(type, prefix, xlate_flags) \  
inline type##_t prefix##_type(reg_t addr, bool  
require_alignment = false) { \  
    if ((xlate_flags) != 0) \  
        flush_tlb(); \  
    if (unlikely(addr & (sizeof(type##_t)-1))) { \  
        if (require_alignment)  
load_reserved_address_misaligned(addr); \  
        else return misaligned_load(addr, sizeof(type##_t)); \  
    } \  
    reg_t vpn = addr >> PGSHIFT; \  
    size_t size = sizeof(type##_t); \  
    if (likely(tlb_load_tag[vpn % TLB_ENTRIES] == vpn)) { \  
        if (proc) READ_MEM(addr, size); \  
        return  
from_target(*(target_endian<type##_t>*)(tlb_data[vpn %  
TLB_ENTRIES].host_offset + addr)); \  
    } \  
    if (unlikely(tlb_load_tag[vpn % TLB_ENTRIES] == (vpn |  
TLB_CHECK_TRIGGERS))) { \  

```

```
type##_t data =  
from_target(*(target_endian<type##_t>*)(  
tlb_data[vpn % TLB_ENTRIES].host_offset + addr)); \  
    if (!matched_trigger) { \  
        matched_trigger =  
trigger_exception(OPERATION_LOAD, addr, data); \  
        if (matched_trigger) \  
            throw *matched_trigger; \  
    } \  
    if (proc) READ_MEM(addr, size); \  
    return data; \  
} \  
target_endian<type##_t> res; \  
load_slow_path(addr, sizeof(type##_t),  
(uint8_t*)&res, (xlate_flags)); \  
    if (proc) READ_MEM(addr, size); \  
    if ((xlate_flags) != 0) \  
        flush_tlb(); \  
    return from_target(res); \  
}
```



# Spike内存系统—加载内存


```
WRITE_RVC_RS2S(MMU.load_int32(RVC_RS1S + insn.rvc_lw_imm()));
```

```
load_func(int16, load, 0)  
load_func(int32, load, 0)
```

```
#define load_func(type, prefix, xlate_flags) \  
inline type##_t prefix##_type(reg_t addr, bool  
require_alignment = false) { \  
    if ((xlate_flags) != 0) \  
        flush_tlb(); \  
    if (unlikely(addr & (sizeof(type##_t)-1))) { \  
        if (require_alignment)  
load_reserved_address_misaligned(addr); \  
        else return misaligned_load(addr, sizeof(type##_t)); \  
    } \  
    reg_t vpn = addr >> PGSHIFT; \  
    size_t size = sizeof(type##_t); \  
    if (likely(tlb_load_tag[vpn % TLB_ENTRIES] == vpn)) { \  
        if (proc) READ_MEM(addr, size); \  
        return  
from_target(*(target_endian<type##_t>*)(tlb_data[vpn %  
TLB_ENTRIES].host_offset + addr)); \  
    } \  
    if (unlikely(tlb_load_tag[vpn % TLB_ENTRIES] == (vpn |  
TLB_CHECK_TRIGGERS))) { \  

```

```
type##_t data =  
from_target(*(target_endian<type##_t>*)(  
tlb_data[vpn % TLB_ENTRIES].host_offset + addr)); \  
    if (!matched_trigger) { \  
        matched_trigger =  
trigger_exception(OPERATION_LOAD, addr, data); \  
        if (matched_trigger) \  
            throw *matched_trigger; \  
    } \  
    if (proc) READ_MEM(addr, size); \  
    return data; \  
} \  
target_endian<type##_t> res; \  
load_slow_path(addr, sizeof(type##_t),  
(uint8_t*)&res, (xlate_flags)); \  
    if (proc) READ_MEM(addr, size); \  
    if ((xlate_flags) != 0) \  
        flush_tlb(); \  
    return from_target(res); \  
}
```




# Spike内存系统—加载数据慢速路径

riscv/mmu.cc

```
void mmu_t::load_slow_path(reg_t addr, reg_t len, uint8_t* bytes, uint32_t
xlate_flags)
{
    reg_t paddr = translate(addr, len, LOAD, xlate_flags);

    if (auto host_addr = sim->addr_to_mem(paddr)) {
        memcpy(bytes, host_addr, len);
        if (tracer.interested_in_range(paddr, paddr + PGSIZE, LOAD))
            tracer.trace(paddr, len, LOAD);
        else
            refill_tlb(addr, paddr, host_addr, LOAD);
    } else if (!mmio_load(paddr, len, bytes)) {
        throw trap_load_access_fault((proc) ? proc->state.v : false, addr, 0, 0);
    }

    if (!matched_trigger) {
        reg_t data = reg_from_bytes(len, bytes);
        matched_trigger = trigger_exception(OPERATION_LOAD, addr, data);
        if (matched_trigger)
            throw *matched_trigger;
    }
}
```




# Spike内存系统—存储内存

```
MMU.store_uint32(RVC_RS1S + insn.rvc_lw_imm(), RVC_RS2S);
```

```
#define store_func(type, prefix, xlate_flags) \
void prefix##_##type(reg_t addr, type##_t val) { \
    if ((xlate_flags) != 0) \
        flush_tlb(); \
    if (unlikely(addr & (sizeof(type##_t)-1))) \
        return misaligned_store(addr, val, sizeof(type##_t)); \
    \
    reg_t vpn = addr >> PGSHIFT; \
    size_t size = sizeof(type##_t); \
    if (likely(tlb_store_tag[vpn % TLB_ENTRIES] == \
        vpn)) { \
        if (proc) WRITE_MEM(addr, val, size); \
        *(target_endian<type##_t>*)(tlb_data[vpn % \
        TLB_ENTRIES].host_offset + addr) = to_target(val); \
    } \
    else if (unlikely(tlb_store_tag[vpn % TLB_ENTRIES] \
        == (vpn | TLB_CHECK_TRIGGERS))) { \
        if (!matched_trigger) { \
            matched_trigger = \
            trigger_exception(OPERATION_STORE, addr, val); \
        } \
    } \
}
```

```
store_func(uint8, store, 0) \
store_func(uint16, store, 0) \
store_func(uint32, store, 0) \
store_func(uint64, store, 0)
```

```
if (matched_trigger) \
    throw *matched_trigger; \
} \
if (proc) WRITE_MEM(addr, val, size); \
*(target_endian<type##_t>*)(tlb_data[vpn % \
    TLB_ENTRIES].host_offset + addr) = to_target(val); \
} \
else { \
    target_endian<type##_t> target_val = \
    to_target(val); \
    store_slow_path(addr, sizeof(type##_t), (const \
    uint8_t*)&target_val, (xlate_flags)); \
    if (proc) WRITE_MEM(addr, val, size); \
} \
if ((xlate_flags) != 0) \
    flush_tlb(); \
}
```




# Spike内存系统—存储内存

```
MMU.store_uint32(RVC_RS1S + insn.rvc_lw_imm(), RVC_RS2S);
```

```
#define store_func(type, prefix, xlate_flags) \
void prefix##_##type(reg_t addr, type##_t val) { \
    if ((xlate_flags) != 0) \
        flush_tlb(); \
    if (unlikely(addr & (sizeof(type##_t)-1))) \
        return misaligned_store(addr, val, sizeof(type##_t)); \
    \
    reg_t vpn = addr >> PGSHIFT; \
    size_t size = sizeof(type##_t); \
    if (likely(tlb_store_tag[vpn % TLB_ENTRIES] == \
        vpn)) { \
        if (proc) WRITE_MEM(addr, val, size); \
        *(target_endian<type##_t>*)(tlb_data[vpn % \
        TLB_ENTRIES].host_offset + addr) = to_target(val); \
    } \
    else if (unlikely(tlb_store_tag[vpn % TLB_ENTRIES] \
        == (vpn | TLB_CHECK_TRIGGERS))) { \
        if (!matched_trigger) { \
            matched_trigger = \
            trigger_exception(OPERATION_STORE, addr, val); \
        } \
    } \
}
```

```
store_func(uint8, store, 0) \
store_func(uint16, store, 0) \
store_func(uint32, store, 0) \
store_func(uint64, store, 0)
```

```
if (matched_trigger) \
    throw *matched_trigger; \
} \
if (proc) WRITE_MEM(addr, val, size); \
*(target_endian<type##_t>*)(tlb_data[vpn % \
    TLB_ENTRIES].host_offset + addr) = to_target(val); \
} \
else { \
    target_endian<type##_t> target_val = \
    to_target(val); \
    store_slow_path(addr, sizeof(type##_t), (const \
    uint8_t*)&target_val, (xlate_flags)); \
    if (proc) WRITE_MEM(addr, val, size); \
} \
if ((xlate_flags) != 0) \
    flush_tlb(); \
}
```




# Spike内存系统—存储慢速路径

riscv/mmu.cc

```
void mmu_t::store_slow_path(reg_t addr, reg_t len, const uint8_t* bytes,
uint32_t xlate_flags)
{
    reg_t paddr = translate(addr, len, STORE, xlate_flags);

    if (!matched_trigger) {
        reg_t data = reg_from_bytes(len, bytes);
        matched_trigger = trigger_exception(OPERATION_STORE, addr, data);
        if (matched_trigger)
            throw *matched_trigger;
    }

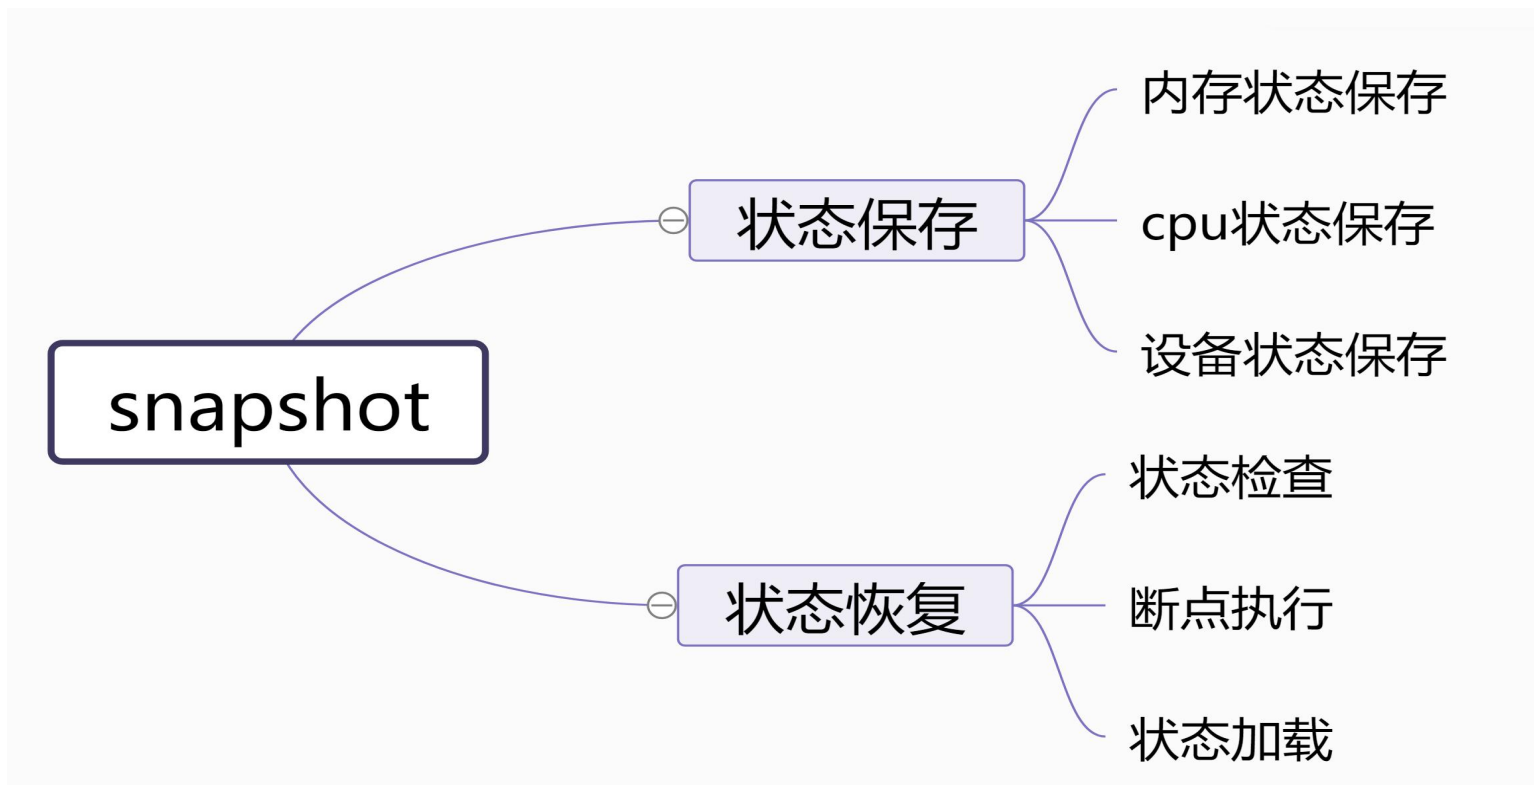
    if (auto host_addr = sim->addr_to_mem(paddr)) {
        memcpy(host_addr, bytes, len);
        if (tracer.interested_in_range(paddr, paddr + PGSIZE, STORE))
            tracer.trace(paddr, len, STORE);
        else
            refill_tlb(addr, paddr, host_addr, STORE);
    } else if (!mmio_store(paddr, len, bytes)) {
        throw trap_store_access_fault((proc) ? proc->state.v : false, addr, 0, 0);
    }
}
```





# snapshot机制回顾

snapshot(快照)功能，可以用来保存系统运行到某一时刻的状态，提供某一时刻系统运行状态的完全拷贝，同时支持从该保存点继续运行的能力



# snapshot机制—写内存保存

## 对写内存进行标记

```
void mmu_t::store_slow_path(reg_t addr, reg_t len, const uint8_t* bytes, uint32_t xlate_flags)
{
    reg_t paddr = translate(addr, len, STORE, xlate_flags);

    if (!matched_trigger) {
        reg_t data = reg_from_bytes(len, bytes);
        matched_trigger = trigger_exception(OPERATION_STORE, addr, data);
        if (matched_trigger)
            throw *matched_trigger;
    }

    if (auto host_addr = sim->addr_to_mem(paddr)) {
        (*sim->get_tags())[paddr >> PGSHIFT] = true;
        memcpy(host_addr, bytes, len);
        if (tracer.interested_in_range(paddr, paddr + PGSIZE, STORE))
            tracer.trace(paddr, len, STORE);
        else
            refill_tlb(addr, paddr, host_addr, STORE);
    } else if (!mmio_store(paddr, len, bytes)) {
        throw trap_store_access_fault(addr, 0, 0);
    }
}
```

# snapshot机制—写内存保存

## 只dump标记为写的内存

snapshot/  
ramdump.cc

```
bool snapshot_t::ramdump(ofstream &out)
{
    for (auto tag: *tags) {
        out.write((char *)&tag.first, sizeof(tag.first)); // vaddr
        size_t size = 1 << PGSHIFT;
        char *addr = sim -> addr_to_mem(tag.first * size);
        if(addr == NULL)
            return false;
        out.write((char *)addr, size);
    }
    return true;
}
```

# 谢谢各位

欢迎提问、讨论、交流合作