

# 从零开始的RISC-V模拟器开发

## 第5讲 Spike篇之设备模拟

中国科学院软件研究所  
PLCT实验室

李威威 liweiwei@iscas.ac.cn

王俊强 wangjunqiang@iscas.ac.cn

吴伟 wuwei2016@iscas.ac.cn

## 回顾


## riscv/mmu.cc

```
tlb_entry_t mmu_t::fetch_slow_path(reg_t vaddr)
{
    reg_t paddr = translate(vaddr, sizeof(fetch_temp), FETCH, 0);

    if (auto host_addr = sim->addr_to_mem(paddr)) {
        return refill_tlb(vaddr, paddr, host_addr, FETCH);
    } else {
        if (!mmio_load(paddr, sizeof fetch_temp, (uint8_t*)&fetch_temp))
            throw trap_instruction_access_fault(vaddr, 0, 0);
        tlb_entry_t entry = {(char*)&fetch_temp - vaddr, paddr - vaddr};
        return entry;
    }
}
```

## riscv/sim.cc

```
char* sim_t::addr_to_mem(reg_t addr) {
    if (!paddr_ok(addr))
        return NULL;
    auto desc = bus.find_device(addr);
    if (auto mem = dynamic_cast<mem_t*>(desc.second))
        if (addr - desc.first < mem->size())
            return mem->contents() + (addr - desc.first);
    return NULL;
}
```



# mmio访问


riscv/mmu.cc

```
bool mmu_t::mmio_load(reg_t addr, size_t len, uint8_t* bytes)
{
    if (!mmio_ok(addr, LOAD))
        return false;

    return sim->mmio_load(addr, len, bytes);
}

bool mmu_t::mmio_store(reg_t addr, size_t len, const uint8_t*
bytes)
{
    if (!mmio_ok(addr, STORE))
        return false;

    return sim->mmio_store(addr, len, bytes);
}
```



riscv/sim.cc

```
bool sim_t::mmio_load(reg_t addr, size_t len, uint8_t* bytes)
{
    if (addr + len < addr || !paddr_ok(addr + len - 1))
        return false;
    return bus.load(addr, len, bytes);
}

bool sim_t::mmio_store(reg_t addr, size_t len, const uint8_t* bytes)
{
    if (addr + len < addr || !paddr_ok(addr + len - 1))
        return false;
    return bus.store(addr, len, bytes);
}
```

# Spike设备树

## 设备树dts

命令: --dump-dts

```
std::string make_dts(size_t insns_per_rtc_tick, size_t cpu_hz,  
    reg_t initrd_start, reg_t initrd_end,  
    const char* bootargs,  
    std::vector<processor_t*> procs,  
    std::vector<std::pair<reg_t, mem_t*>> mems)
```

```
liww@liww-tm: /workspace/riscv/plct-spike/build
文件(F) 编辑(E) 查看(V) 搜索(S) 终端(T) 帮助(H)

cpus {
    #address-cells = <1>;
    #size-cells = <0>;
    timebase-frequency = <100000000>;
    CPU0: cpu@0 {
        device_type = "cpu";
        reg = <0>;
        status = "okay";
        compatible = "riscv";
        riscv,isa = "rv64imafdc";
        mmu-type = "riscv,sv48";
        riscv,pmpregions = <16>;
        riscv,pmpgranularity = <4>;
        clock-frequency = <1000000000>;
        CPU0_intc: interrupt-controller {
            #interrupt-cells = <1>;
            interrupt-controller;
            compatible = "riscv,cpu-intc";
        };
    };
};

memory@80000000 {
    device_type = "memory";
    reg = <0x0 0x80000000 0x0 0x80000000>;
};

soc {
    #address-cells = <2>;
    #size-cells = <2>;
    compatible = "ucbbar,spike-bare-soc", "simple-bus";
    ranges;
    clint@2000000 {
        compatible = "riscv,clint0";
        interrupts-extended = <&CPU0_intc 3 &CPU0_intc 7 >;
        reg = <0x0 0x2000000 0x0 0xc0000>;
    };
};

htif {
    compatible = "ucb,htif0";
};
};

liww@liww-tm:build$
```

# Spike设备支持

设备基础类型

riscv/devices.h

```
class abstract_device_t {  
public:  
    virtual bool load(reg_t addr, size_t len, uint8_t* bytes) = 0;  
    virtual bool store(reg_t addr, size_t len, const uint8_t* bytes) =  
0;  
    virtual ~abstract_device_t() {}  
};
```

# Spike设备支持

设备基础类型

riscv/devices.h

```
class abstract_device_t {  
    public:  
        virtual bool load(reg_t addr, size_t len, uint8_t* bytes) = 0;  
        virtual bool store(reg_t addr, size_t len, const uint8_t* bytes) =  
0;  
        virtual ~abstract_device_t() {}  
};
```

```
class processor_t : public  
abstract_device_t  
{  
    ...  
};
```

# Spike设备支持

设备基础类型

riscv/devices.h

```
class abstract_device_t {  
public:  
    virtual bool load(reg_t addr, size_t len, uint8_t* bytes) = 0;  
    virtual bool store(reg_t addr, size_t len, const uint8_t* bytes) =  
0;  
    virtual ~abstract_device_t() {}  
};
```

bus

mem

rom

clint

plugin

# Spike设备支持

riscv/sim.cc: sim\_t::sim\_t(...)

```
for (auto& x : mems)
    bus.add_device(x.first, x.second);

for (auto& x : plugin_devices)
    bus.add_device(x.first, x.second);
...
//handle clic
clint.reset(new clint_t(procs, CPU_HZ / INSNS_PER_RTC_TICK,
real_time_clint));
reg_t clint_base;
if (fdt_parse_clint(fdt, &clint_base, "riscv,clint0")) {
    bus.add_device(CLINT_BASE, clint.get());
} else {
    bus.add_device(clint_base, clint.get());
}
```

```
void sim_t::set_rom()
{
    ...
    boot_rom.reset(new rom_device_t(rom));
    bus.add_device(DEFAULT_RSTVEC, boot_rom.get());
}
```



# bus设备支持

riscv/devices.cc

riscv/devices.h

```
class bus_t : public abstract_device_t {  
public:  
    bool load(reg_t addr, size_t len, uint8_t* bytes);  
    bool store(reg_t addr, size_t len, const uint8_t* bytes);  
    void add_device(reg_t addr, abstract_device_t* dev);  
  
    std::pair<reg_t, abstract_device_t*> find_device(reg_t addr);  
  
private:  
    std::map<reg_t, abstract_device_t*> devices;  
};
```

```
void bus_t::add_device(reg_t addr, abstract_device_t* dev)  
{  
    devices[addr] = dev;  
}
```

```
bool bus_t::load(reg_t addr, size_t len, uint8_t* bytes)  
{  
    // Find the device with the base address closest to but  
    // less than addr (price-is-right search)  
    auto it = devices.upper_bound(addr);  
    if (devices.empty() || it == devices.begin()) {  
        // Either the bus is empty, or there weren't  
        // any items with a base address <= addr  
        return false;  
    }  
    // Found at least one item with base address <= addr  
    // The iterator points to the device after this, so  
    // go back by one item.  
    it--;  
    return it->second->load(addr - it->first, len, bytes);  
}
```

# mem设备支持

```
class mem_t : public abstract_device_t {  
public:  
    mem_t(reg_t size);  
    mem_t(const mem_t& that) = delete;  
    ~mem_t();
```

```
    bool load(reg_t addr, size_t len, uint8_t* bytes) { return load_store(addr, len, bytes, false); }  
    bool store(reg_t addr, size_t len, const uint8_t* bytes) { return load_store(addr, len, const_cast<uint8_t*>(bytes),  
true); }  
    char* contents(reg_t addr);  
    reg_t size() { return SZ; }
```

```
private:  
    bool load_store(reg_t addr, size_t len, uint8_t* bytes, bool store);
```

```
    std::map<reg_t, char*> sparse_memory_map;  
    reg_t SZ;  
};
```

```
mem_t::mem_t(reg_t size)  
    : SZ(size)  
    {  
        if (size == 0 || size % PGSIZE != 0)  
            throw std::runtime_error("memory size must be a positive multiple of 4  
KiB");  
    }
```

# mem设备支持

spike\_main/spike.cc

```
parser.option('m', 0, 1, [&](const char* s){mems =  
make_mems(s);});  
...  
if (mems.empty())  
    mems = make_mems("2048");
```



```
for (auto& x : mems)  
    bus.add_device(x.first, x.second);
```



```
char* sim_t::addr_to_mem(reg_t addr) {  
    if (!paddr_ok(addr))  
        return NULL;  
    auto desc = bus.find_device(addr);  
    if (auto mem = dynamic_cast<mem_t*>(desc.second))  
        if (addr - desc.first < mem->size())  
            return mem->contents() + (addr - desc.first);  
    return NULL;  
}
```

命令	描述
-m<n>	Provide <n> MiB of target memory [default 2048]
-m<a:m, b:n,...>	Provide memory regions of size m and n bytes at base addresses a and b (with 4 KiB alignment)

# mem设备支持

```
char* mem_t::contents(reg_t addr) {  
    reg_t ppn = addr >> PGSHIFT, pgoff = addr % PGSIZE;  
    auto search = sparse_memory_map.find(ppn);  
    if (search == sparse_memory_map.end()) {  
        auto res = (char*)calloc(PGSIZE, 1);  
        if (res == nullptr)  
            throw std::bad_alloc();  
        sparse_memory_map[ppn] = res;  
        return res + pgoff;  
    }  
    return search->second + pgoff;  
}
```

# rom设备支持

## riscv/devices.h

```
class rom_device_t : public abstract_device_t {  
public:  
    rom_device_t(std::vector<char> data);  
    bool load(reg_t addr, size_t len, uint8_t* bytes);  
    bool store(reg_t addr, size_t len, const uint8_t* bytes);  
    const std::vector<char>& contents() { return data; }  
private:  
    std::vector<char> data;  
};
```

## riscv/rom.cc

```
rom_device_t::rom_device_t(std::vector<char> data)  
    : data(data)  
{  
}  
  
bool rom_device_t::load(reg_t addr, size_t len, uint8_t* bytes)  
{  
    if (addr + len > data.size())  
        return false;  
    memcpy(bytes, &data[addr], len);  
    return true;  
}  
  
bool rom_device_t::store(reg_t addr, size_t len, const uint8_t* bytes)  
{  
    return false;  
}
```

# clint设备支持

```
class clint_t : public abstract_device_t {
public:
    clint_t(std::vector<processor_t*>&, uint64_t freq_hz, bool real_time);
    bool load(reg_t addr, size_t len, uint8_t* bytes);
    bool store(reg_t addr, size_t len, const uint8_t* bytes);
    size_t size() { return CLINT_SIZE; }
    void increment(reg_t inc);
private:
    typedef uint64_t mtime_t;
    typedef uint64_t mtimecmp_t;
    typedef uint32_t msip_t;
    std::vector<processor_t*>& procs;
    uint64_t freq_hz;
    bool real_time;
    uint64_t real_time_ref_secs;
    uint64_t real_time_ref_usecs;
    mtime_t mtime;
    std::vector<mtimecmp_t> mtimecmp;
};
```

```
/* 0000 msip hart 0
 * 0004 msip hart 1
 * 4000 mtimecmp hart 0 lo
 * 4004 mtimecmp hart 0 hi
 * 4008 mtimecmp hart 1 lo
 * 400c mtimecmp hart 1 hi
 * bff8 mtime lo
 * bffc mtime hi
 */
```

```
clint_t::clint_t(std::vector<processor_t*>& procs, uint64_t
freq_hz, bool real_time)
: procs(procs), freq_hz(freq_hz), real_time(real_time), mtime(0),
mtimecmp(procs.size())
{
    struct timeval base;

    gettimeofday(&base, NULL);

    real_time_ref_secs = base.tv_sec;
    real_time_ref_usecs = base.tv_usec;
}
```

# clint设备支持

```
parser.option(0, "real-time-clint", 0, [&](const char *s){real_time_clint = true;});
```

riscv/sim.cc: sim\_t::sim\_t(...)

```
clint.reset(new clint_t(procs, CPU_HZ / INSNS_PER_RTC_TICK, real_time_clint));  
reg_t clint_base;  
if (fdt_parse_clint(fdt, &clint_base, "riscv,clint0")) {  
    bus.add_device(CLINT_BASE, clint.get());  
} else {  
    bus.add_device(clint_base, clint.get());  
}
```

# clint设备支持

```
void sim_t::step(size_t n)
{
    for (size_t i = 0, steps = 0; i < n; i += steps)
    {
        steps = std::min(n - i, INTERLEAVE - current_step);
        procs[current_proc]->step(steps);

        current_step += steps;
        if (current_step == INTERLEAVE)
        {
            current_step = 0;
            procs[current_proc]->get_mmu()->yield_load_reservation();
            if (++current_proc == procs.size()) {
                current_proc = 0;
                clint->increment(INTERLEAVE / INSNS_PER_RTC_TICK);
            }

            host->switch_to();
        }
    }
}
```

```
void clint_t::increment(reg_t inc)
{
    if (real_time) {
        struct timeval now;
        uint64_t diff_usecs;

        gettimeofday(&now, NULL);
        diff_usecs = ((now.tv_sec - real_time_ref_secs) * 1000000) + (now.tv_usec
- real_time_ref_usecs);
        mtime = diff_usecs * freq_hz / 1000000;
    } else {
        mtime += inc;
    }
    for (size_t i = 0; i < procs.size(); i++) {
        procs[i]->state.mip &= ~MIP_MTIP;
        if (mtime >= mtimecmp[i])
            procs[i]->state.mip |= MIP_MTIP;
    }
}
```



# Spike设备扩展

Mmio plugin设备

riscv/devices.cc

riscv/devices.h

```
class mmio_plugin_device_t : public abstract_device_t {  
public:  
    mmio_plugin_device_t(const std::string& name, const  
std::string& args);  
    virtual ~mmio_plugin_device_t() override;  
  
    virtual bool load(reg_t addr, size_t len, uint8_t* bytes) override;  
    virtual bool store(reg_t addr, size_t len, const uint8_t* bytes)  
override;  
  
private:  
    mmio_plugin_t plugin;  
    void* user_data;  
};
```

```
mmio_plugin_device_t::mmio_plugin_device_t(const std::string& name,  
                                             const std::string& args)  
    : plugin(mmio_plugin_map().at(name)),  
      user_data((*plugin.alloc)(args.c_str()))  
{  
}  
  
mmio_plugin_device_t::~~mmio_plugin_device_t()  
{  
    (*plugin.dealloc)(user_data);  
}  
  
bool mmio_plugin_device_t::load(reg_t addr, size_t len, uint8_t* bytes)  
{  
    return (*plugin.load)(user_data, addr, len, bytes);  
}  
  
bool mmio_plugin_device_t::store(reg_t addr, size_t len, const uint8_t*  
bytes)  
{  
    return (*plugin.store)(user_data, addr, len, bytes);  
}
```

# Spike plugin设备接口

riscv/mmio\_plugin.h

```
typedef struct {  
    // Allocate user data for an instance of the plugin. The parameter is a simple  
    // c-string containing arguments used to construct the plugin. It returns a  
    // void* to the allocated data.  
    void* (*alloc)(const char*);  
  
    // Load a memory address of the MMIO plugin. The parameters are the user_data  
    // (void*), memory offset (reg_t), number of bytes to load (size_t), and the  
    // buffer into which the loaded data should be written (uint8_t*). Return true  
    // if the load is successful and false otherwise.  
    bool (*load)(void*, reg_t, size_t, uint8_t*);  
  
    // Store some bytes to a memory address of the MMIO plugin. The parameters are  
    // the user_data (void*), memory offset (reg_t), number of bytes to store  
    // (size_t), and the buffer containing the data to be stored (const uint8_t*).  
    // Return true if the store is successful and false otherwise.  
    bool (*store)(void*, reg_t, size_t, const uint8_t*);  
  
    // Deallocate the data allocated during the call to alloc. The parameter is a  
    // pointer to the user data allocated during the call to alloc.  
    void (*dealloc)(void*);  
} mmio_plugin_t;
```

# Spike plugin设备注册

riscv/devices.cc

```
// Type for holding all registered MMIO plugins by name.
using mmio_plugin_map_t = std::map<std::string,
mmio_plugin_t>;

void register_mmio_plugin(const char* name_cstr,
                          const mmio_plugin_t* mmio_plugin)
{
    std::string name(name_cstr);
    if (!mmio_plugin_map().emplace(name, *mmio_plugin).second) {
        throw std::runtime_error("Plugin \"" + name + "\" already
registered!");
    }
}
```

# Spike设备扩展

## 以块设备为例

```
__attribute__((constructor)) static void on_load()
{
    static mmio_plugin_t test_mmio_plugin = {
        test_mmio_plugin_alloc,
        test_mmio_plugin_load,
        test_mmio_plugin_store,
        test_mmio_plugin_dealloc
    };

    register_mmio_plugin("test_mmio_plugin",
        &test_mmio_plugin);
}
```

```
{
    vblock@10000000 {
        compatible = "plct,vblock";
        reg = <0x0 0x10000000 0x0 0xa00000>;
        len = <0x5000>;
    };
}
```

代码仓库: <https://github.com/plctlab/plct-spike/tree/snapshot-develop>

# Spike设备扩展

## 以块设备为例

```
void* test_mmio_plugin_alloc(const char* args)
{
    printf("ALLOC -- ARGS=%s\n", args);
    int fd = open(args, O_RDWR, 0);
    if(fd==-1){
        printf("open block disk file failed.\n");
        exit(-1);
    }
    struct stat st; //定义文件信息结构体
    int r=fstat(fd,&st);
    if(r==-1){
        printf("get file size failed. \n");
        close(fd);
        exit(-1);
    }
    len=st.st_size;
    void *
    p=mmap(NULL,len,PROT_READ|PROT_WRITE,MAP_SHARED,fd,0);
    return p;
}
```

```
bool test_mmio_plugin_load(void* self, reg_t addr, size_t len,
uint8_t* bytes)
{
    memcpy(bytes, (char *)self + addr, len);
    return true;
}

bool test_mmio_plugin_store(void* self, reg_t addr, size_t len,
const uint8_t* bytes)
{
    memcpy((char *)self + addr, bytes, len);
    msync((char *)self + addr,len,0);
    return true;
}

void test_mmio_plugin_dealloc(void* self)
{
    munmap(self, len);
    printf("DEALLOC -- SELF=%p\n", self);
}
```

# HTIF

## Host/Target Interface

- riscv-fesvr runs on your host computer and communicates with the target design
- a non-standard tool for Berkeley processors

```
class sim_t : public htif_t, public simif_t
{
    ...
}
```

```
int sim_t::run()
{
    host = context_t::current();
    target.init(sim_thread_main, this);
    return htif_t::run();
}
```

# HTIF初始化

```
htif_t::htif_t(int argc, char** argv) : htif_t()
{
    parse_arguments(argc, argv);
    register_devices();
}
```

# HTIF参数解析

```
htif_t::htif_t(int argc, char** argv) : htif_t()
{
    parse_arguments(argc, argv);
    register_devices();
}
```



```
void htif_t::parse_arguments(int argc, char ** argv)
{
    ...
    if (c == -1) break;
retry:
    switch (c) {
        case 'h': usage(argv[0]);
            throw std::invalid_argument("User queried htif_t help text");
        case HTIF_LONG_OPTIONS_OPTIND:
            if (optarg) dynamic_devices.push_back(new rfb_t(atoi(optarg)));
            else dynamic_devices.push_back(new rfb_t);
            break;
        case HTIF_LONG_OPTIONS_OPTIND + 1:
            // [TODO] Remove once disks are supported again
            throw std::invalid_argument("--disk/+disk unsupported (use a ramdisk)");
            dynamic_devices.push_back(new disk_t(optarg));
            Break;
        ...
    }
}
```




# HTIF选项

```
#define HTIF_USAGE_OPTIONS \  
"HOST OPTIONS\n\  
-h, --help      Display this help and exit\n\  
+h, +help\n    +permissive      The host will ignore any unparsed options up until\n    +permissive-off  (Only needed for VCS)\n    +permissive-off  Stop ignoring options. This is mandatory if using\n    +permissive      (Only needed for VCS)\n--rfb=DISPLAY    Add new remote frame buffer on display DISPLAY\n+rfb=DISPLAY      to be accessible on 5900 + DISPLAY (default = 0)\n--signature=FILE  Write torture test signature to FILE\n+signature=FILE\n--signature-granularity=VAL    Size of each line in signature.\n+signature-granularity=VAL\n--chroot=PATH      Use PATH as location of syscall-servicing binaries\n+chroot=PATH\n--payload=PATH      Load PATH memory as an additional ELF payload\n+payload=PATH\n\n\  
HOST OPTIONS (currently unsupported)\n--disk=DISK        Add DISK device. Use a ramdisk since this isn't\n+disk=DISK          supported\n\n\  

```

# HTIF设备注册

```
htif_t::htif_t(int argc, char** argv) : htif_t()
{
    parse_arguments(argc, argv);
    register_devices();
}
```



```
void htif_t::register_devices()
{
    device_list.register_device(&syscall_proxy);
    device_list.register_device(&bcd);
    for (auto d : dynamic_devices)
        device_list.register_device(d);
}
```

```
device_list_t device_list;
syscall_t syscall_proxy;
bcd_t bcd;
std::vector<device_t*> dynamic_devices;
```

```
void htif_t::parse_arguments(int argc, char** argv)
{
    ...
    if (c == -1) break;
retry:
    switch (c) {
        case 'h': usage(argv[0]);
            throw std::invalid_argument("User queried htif_t help text");
        case HTIF_LONG_OPTIONS_OPTIND:
            if (optarg) dynamic_devices.push_back(new rfb_t(atoi(optarg)));
            else dynamic_devices.push_back(new rfb_t);
            break;
        case HTIF_LONG_OPTIONS_OPTIND + 1:
            // [TODO] Remove once disks are supported again
            throw std::invalid_argument("--disk/+disk unsupported (use a ramdisk)");
            dynamic_devices.push_back(new disk_t(optarg));
            Break;
        ...
    }
}
```

# device list

```
class device_list_t
{
public:
    device_list_t();
    void register_device(device_t* dev);
    void handle_command(command_t cmd);
    void tick();

private:
    std::vector<device_t*> devices;
    null_device_t null_device;
    size_t num_devices;
};
```

```
void device_list_t::register_device(device_t* dev)
{
    num_devices++;
    assert(num_devices < command_t::MAX_DEVICES);
    devices[num_devices-1] = dev;
}

void device_list_t::handle_command(command_t cmd)
{
    devices[cmd.device()->handle_command(cmd);
}

void device_list_t::tick()
{
    for (size_t i = 0; i < num_devices; i++)
        devices[i]->tick();
}
```

# HTIF运行

```
int htif_t::run()
{
    start();

    auto enq_func = [](std::queue<reg_t>* q, uint64_t x) { q->push(x); };
    std::queue<reg_t> fromhost_queue;
    std::function<void(reg_t)> fromhost_callback =
    std::bind(enq_func, &fromhost_queue, std::placeholders::_1);

    if (tohost_addr == 0) {
        while (true)
            idle();
    }

    while (!signal_exit && exitcode == 0)
    {
        if (auto tohost = from_target(mem.read_uint64(tohost_addr))) {
            mem.write_uint64(tohost_addr, target_endian<uint64_t>::zero);
            command_t cmd(mem, tohost, fromhost_callback);
            device_list.handle_command(cmd);
        } else {
            idle();
        }
    }
}
```

```
device_list.tick();

if (!fromhost_queue.empty() && !mem.read_uint64(fromhost_addr)) {
    mem.write_uint64(fromhost_addr, to_target(fromhost_queue.front()));
    fromhost_queue.pop();
}

stop();

return exit_code();
}
```

# HTIF运行

```
int htif_t::run()
{
    start();

    auto enq_func = [](std::queue<reg_t>* q, uint64_t x) { q->push(x); };
    std::queue<reg_t> fromhost_queue;
    std::function<void(reg_t)> fromhost_callback =
    std::bind(enq_func, &fromhost_queue, std::placeholders::_1);

    if (tohost_addr == 0) {
        while (true)
            idle();
    }

    while (!signal_exit && exitcode == 0)
    {
        if (auto tohost = from_target(mem.read_uint64(tohost_addr))) {
            mem.write_uint64(tohost_addr, target_endian<uint64_t>::zero);
            command_t cmd(mem, tohost, fromhost_callback);
            device_list.handle_command(cmd);
        } else {
            idle();
        }
    }
}
```

```
device_list.tick();
```

```
if (!fromhost_queue.empty() && !mem.read_uint64(fromhost_addr)) {
    mem.write_uint64(fromhost_addr, to_target(fromhost_queue.front()));
    fromhost_queue.pop();
}
```

```
stop();
```

```
return exit_code();
}
```

# HTIF设备处理及注册命令

```
void device_list_t::handle_command(command_t cmd)
{
    devices[cmd.device()->handle_command(cmd);
}
```



```
void device_t::handle_command(command_t cmd)
{
    command_handlers[cmd.cmd()](cmd);
}
```



```
void device_t::register_command(size_t cmd, command_func_t handler, const char* name)
{
    assert(cmd < command_t::MAX_COMMANDS);
    assert(strlen(name) < IDENTITY_SIZE);
    command_handlers[cmd] = handler;
    command_names[cmd] = name;
}
```

# HTIF设备结构

```
class device_t
{
public:
    device_t();
    virtual ~device_t() {}
    virtual const char* identity() = 0;
    virtual void tick() {}

    void handle_command(command_t cmd);

protected:
    typedef std::function<void(command_t)> command_func_t;
    void register_command(size_t, command_func_t, const char*);

private:
    device_t& operator = (const device_t&); // disallow
    device_t(const device_t&); // disallow
    static const size_t IDENTITY_SIZE = 64;
    void handle_null_command(command_t cmd);
    void handle_identify(command_t cmd);

    std::vector<command_func_t> command_handlers;
    std::vector<std::string> command_names;
};
```

fesvr/device.h

# HTIF命令结构

```
class device_t
{
public:
    device_t();
    virtual ~device_t() {}
    virtual const char* identity() = 0;
    virtual void tick() {}

    void handle_command(command_t cmd);

protected:
    typedef std::function<void(command_t)> command_func_t;
    void register_command(size_t, command_func_t, const char*);

private:
    device_t& operator = (const device_t&); // disallow
    device_t(const device_t&); // disallow
    static const size_t IDENTITY_SIZE = 64;
    void handle_null_command(command_t cmd);
    void handle_identify(command_t cmd);

    std::vector<command_func_t> command_handlers;
    std::vector<std::string> command_names;
};
```

## fesvr/device.h

```
class command_t
{
public:
    typedef std::function<void(uint64_t)> callback_t;
    command_t(memif_t& memif, uint64_t tohost, callback_t cb)
        : _memif(memif), tohost(tohost), cb(cb) {}

    memif_t& memif() { return _memif; }
    uint8_t device() { return tohost >> 56; }
    uint8_t cmd() { return tohost >> 48; }
    uint64_t payload() { return tohost << 16 >> 16; }
    void respond(uint64_t resp) { cb((tohost >> 48 << 48) | (resp << 16 >> 16)); }

    static const size_t MAX_COMMANDS = 256;
    static const size_t MAX_DEVICES = 256;

private:
    memif_t& _memif;
    uint64_t tohost;
    callback_t cb;
};
```



# HTIF系统调用代理设备

fesvr/syscall.h

```
class syscall_t : public device_t
{
public:
    syscall_t(htif_t*);

    void set_chroot(const char* where);

private:
    const char* identity() { return "syscall_proxy"; }

    htif_t* htif;
    memif_t* memif;
    std::vector<syscall_func_t> table;
    fds_t fds;

    void handle_syscall(command_t cmd);
    void dispatch(addr_t mm);

    std::string chroot;
    std::string do_chroot(const char* fn);
    std::string undo_chroot(const char* fn);
```

```
reg_t sys_exit(reg_t, reg_t, reg_t, reg_t, reg_t, reg_t, reg_t);
reg_t sys_openat(reg_t, reg_t, reg_t, reg_t, reg_t, reg_t, reg_t);
reg_t sys_read(reg_t, reg_t, reg_t, reg_t, reg_t, reg_t, reg_t);
reg_t sys_pread(reg_t, reg_t, reg_t, reg_t, reg_t, reg_t, reg_t);
reg_t sys_write(reg_t, reg_t, reg_t, reg_t, reg_t, reg_t, reg_t);
reg_t sys_pwrite(reg_t, reg_t, reg_t, reg_t, reg_t, reg_t, reg_t);
reg_t sys_close(reg_t, reg_t, reg_t, reg_t, reg_t, reg_t, reg_t);
reg_t sys_lseek(reg_t, reg_t, reg_t, reg_t, reg_t, reg_t, reg_t);
reg_t sys_fstat(reg_t, reg_t, reg_t, reg_t, reg_t, reg_t, reg_t);
reg_t sys_lstat(reg_t, reg_t, reg_t, reg_t, reg_t, reg_t, reg_t);
reg_t sys_fstatat(reg_t, reg_t, reg_t, reg_t, reg_t, reg_t, reg_t);
reg_t sys_faccessat(reg_t, reg_t, reg_t, reg_t, reg_t, reg_t, reg_t);
reg_t sys_fcntl(reg_t, reg_t, reg_t, reg_t, reg_t, reg_t, reg_t);
reg_t sys_ftruncate(reg_t, reg_t, reg_t, reg_t, reg_t, reg_t, reg_t);
reg_t sys_renameat(reg_t, reg_t, reg_t, reg_t, reg_t, reg_t, reg_t);
reg_t sys_linkat(reg_t, reg_t, reg_t, reg_t, reg_t, reg_t, reg_t);
reg_t sys_unlinkat(reg_t, reg_t, reg_t, reg_t, reg_t, reg_t, reg_t);
reg_t sys_mkdirat(reg_t, reg_t, reg_t, reg_t, reg_t, reg_t, reg_t);
reg_t sys_getcwd(reg_t, reg_t, reg_t, reg_t, reg_t, reg_t, reg_t);
reg_t sys_getmainvars(reg_t, reg_t, reg_t, reg_t, reg_t, reg_t, reg_t);
reg_t sys_chdir(reg_t, reg_t, reg_t, reg_t, reg_t, reg_t, reg_t);
};
```

# HTIF系统调用代理设备功能

```
syscall_t::syscall_t(htif_t* htif)
: htif(htif), memif(&htif->memif()), table(2048)
{
    table[17] = &syscall_t::sys_getcwd;
    table[25] = &syscall_t::sys_fcntl;
    table[34] = &syscall_t::sys_mkdirat;
    table[35] = &syscall_t::sys_unlinkat;
    table[37] = &syscall_t::sys_linkat;
    table[38] = &syscall_t::sys_renameat;
    table[46] = &syscall_t::sys_ftruncate;
    table[48] = &syscall_t::sys_faccessat;
    table[49] = &syscall_t::sys_chdir;
    table[56] = &syscall_t::sys_openat;
    table[57] = &syscall_t::sys_close;
    table[62] = &syscall_t::sys_lseek;
    table[63] = &syscall_t::sys_read;
    table[64] = &syscall_t::sys_write;
    table[67] = &syscall_t::sys_pread;
    table[68] = &syscall_t::sys_pwrite;
    ...
    register_command(0, std::bind(&syscall_t::handle_syscall,
    this, _1), "syscall");
    ...
}
```

## fesvr/syscall.cc

```
void syscall_t::handle_syscall(command_t cmd)
{
    if (cmd.payload() & 1) // test pass/fail
    {
        htif->exitcode = cmd.payload();
        if (htif->exit_code())
            std::cerr << "*** FAILED *** (tohost = " << htif->exit_code() << ")" << std::endl;
        return;
    }
    else // proxied system call
        dispatch(cmd.payload());

    cmd.respond(1);
}
```

# htif bcd设备

## fesvr/device.h

```
class bcd_t : public device_t
{
public:
    bcd_t();
    const char* identity() { return "bcd"; }
    void tick();

private:
    void handle_read(command_t cmd);
    void handle_write(command_t cmd);

    std::queue<command_t> pending_reads;
};
```

```
bcd_t::bcd_t()
{
    register_command(0, std::bind(&bcd_t::handle_read, this, _1), "read");
    register_command(1, std::bind(&bcd_t::handle_write, this, _1), "write");
}

void bcd_t::handle_read(command_t cmd)
{
    pending_reads.push(cmd);
}

void bcd_t::handle_write(command_t cmd)
{
    canonical_terminal_t::write(cmd.payload());
}

void bcd_t::tick()
{
    int ch;
    if (!pending_reads.empty() && (ch = canonical_terminal_t::read()) != -1)
    {
        pending_reads.front().respond(0x100 | ch);
        pending_reads.pop();
    }
}
```

# 谢谢各位

欢迎提问、讨论、交流合作